

CM30225 Parallel Computing Coursework 2

Lei Zeng

01 January 2015

Contents

1	How to use	1
2	Code Description	2
2.1	Basic idea	2
2.2	Job Splitting	2
2.3	End Calculation	2
2.4	Pseudo code	3
3	Correctness Testing	4
3.1	Testing basic calculation	4
3.2	Further Testing	5
3.3	Parallelism Testing	5
3.3.1	Testing job splitting	5
3.3.2	Testing Averaging	6
4	Scalability Investigation	9
4.1	Time spent	9
4.2	Speedup	10
4.3	Efficiency	10
4.4	Compare with shared memory	10

Chapter 1

How to use

This program takes 2 inputs:

1. dimension: Type: integer, the size of square array.
2. precision: Type: double, the accuracy you require from the final result.

e.g.:

```
./cw2 10 0.1
```

This will randomly generate a 10 x 10 square array and keep iterating until all new values differ from their old values by less than the precision 0.1.

This program will take use of all cores in each node. Specify numbers of cores and nodes you want to use in the job script.

Chapter 2

Code Description

2.1 Basic idea

All processors run the same code except *Root* processor. *Root* ($rank_0$, in my program) broadcasts the square array to the rest processors. Then all processors average their part of rows with regard to rank number. After average, the rest processors send the results of their rows back to *Root*. *Root* receives all new values from other processors and replaces the square array with new values.

2.2 Job Splitting

In an $n * n$ square array with i processors, we have: $avg = \frac{n-2}{i}$ and $extra = (n-2) \bmod i$. All processors with $rank < extra$ will averaging $avg + 1$ rows, and the rest will averaging avg rows.

E.g. a $15 * 15$ square array with 5 processors:

$avg = \frac{15-2}{5} = 2$ and $extra = (15-2) \bmod 5 = 3$, so from $rank_0$ to $rank_4$, they will average 3, 3, 3, 2, 2 rows.

2.3 End Calculation

There is a flag *isEnd*. Before replacing old values with new one, set *isEnd* to *true*. During replacing, if $|old - new| > precision$ occurs, *isEnd* will be set to *false*. At the end of replacing, *Root* broadcasts *isEnd*. If *isEnd* is not *true*, the averaging step will continue.

2.4 Pseudo code

```
main()
{
    init MPI;

    calculate rows to average with regard to rank;

    if (rank == ROOT) init array;

    while (!isEnd)
    {
        MPI_Bcast array;

        averaging;

        if (rank == ROOT)
        {
            MPI_Recv;
            repacing array;
        }
        else
        {
            MPI_Send;
        }

        MPI_Barrier;

        MPI_Bcast isEnd;
    }
}
```

Chapter 3

Correctness Testing

3.1 Testing basic calculation

Test program in single processor to check the results of each calculation with regard to specification. The following sample using 4 x 4 square array with *precision* = 1.0, which is the same as the one I used in *coursework 1*, so it is easier to check results:

84.018772	39.438293	78.309922	79.844003
91.164736	19.755137	33.522276	76.822959
27.777471	55.396996	47.739705	62.887092
36.478447	51.340091	95.222973	91.619507

Each calculation will average 4 neighbours (the old values of top-right-bottom-left cells):

84.018772	39.438293	78.309922	79.844003
91.164736	54.880575	33.522276	76.822959
27.777471	55.396996	47.739705	62.887092
36.478447	51.340091	95.222973	91.619507

Each step will update all cells except boundaries. The result of first update:

84.018772	39.438293	78.309922	79.844003
91.164736	54.880575	55.656931	76.822959
27.777471	36.653101	61.757334	62.887092
36.478447	51.340091	95.222973	91.619507

The difference between new value and old value of cell[2][2] is $54.880575 - 19.755137 = 35.125438$ which is larger than precision (1.0), so the averaging step will be continued. In

this example, it takes 6 steps to end: the new values of all cells differ from their previous values by less than the precision:

84.018772	39.438293	78.309922	79.844003
91.164736	63.671773	72.311806	76.822959
27.777471	53.307976	70.548532	62.887092
36.478447	51.340091	95.222973	91.619507

3.2 Further Testing

Due to the large outputs, it is very difficult to test when the dimension increases. I try to compare the result of each step with my friend who uses different method to accomplish this coursework. I randomly generate a square array with 100 x 100 dimension, pass the file to my friend. We both write the results of each step to an output file. Here is how I compare the output file:

Using git to compare results: Firstly I push my result file to a git repository, and then push my friend's output file as update version of my result. Git will help us to check those two files and highlight the differences.

It turns out that two results are exactly the same, so I will say that both of our programs work as what we expected.

3.3 Parallelism Testing

Here I use another sample to test multi-threading. In my program, the algorithm for splitting work is described as following:

In an $n \times n$ square array with i threads, each thread (except the last one) deals with $\frac{n-2}{i}$ row(s), and the last thread will deal with the rest.

Each thread use the same algorithm which I have tested in 3.1 and 3.2, so I will focus on testing the part of splitting job and parallelism.

3.3.1 Testing job splitting

The out put of testing different dimensions and different number of processors:

1. 20 x 20 with 4 threads:

```
rank 0 in node006: 1 – 5
rank 3 in node006: 15 – 18
rank 1 in node006: 6 – 10
rank 2 in node006: 11 – 14
```

2. 50 x 50 with 8 threads:

```
rank 3 in node006: 19 – 24
rank 1 in node006: 7 – 12
rank 0 in node006: 1 – 6
rank 6 in node006: 37 – 42
rank 5 in node006: 31 – 36
rank 2 in node006: 13 – 18
rank 4 in node006: 25 – 30
rank 7 in node006: 43 – 48
```

3. 100 x 100 with 16 threads:

```
rank 3 in node007: 21 – 26
rank 7 in node007: 45 – 50
rank 0 in node007: 1 – 7
rank 1 in node007: 8 – 14
rank 6 in node007: 39 – 44
rank 5 in node007: 33 – 38
rank 14 in node008: 87 – 92
rank 2 in node007: 15 – 20
rank 4 in node007: 27 – 32
rank 9 in node008: 57 – 62
rank 8 in node008: 51 – 56
rank 10 in node008: 63 – 68
rank 12 in node008: 75 – 80
rank 15 in node008: 93 – 98
rank 11 in node008: 69 – 74
rank 13 in node008: 81 – 86
```

All test cases give the result as expected. The splitting part should work fine.

3.3.2 Testing Averaging

Firstly, I will test a small array: 6 x 6 with 3 processors, so $rank_0$ computes row 1-2, $rank_1$ computes row 3 and $rank_2$ compute row 4.

84.018772	39.438293	78.309922	79.844003	91.164736	19.755137
33.522276	76.822959	27.777471	55.396996	47.739705	62.887092
36.478447	51.340091	95.222973	91.619507	63.571173	71.729693
14.160256	60.696888	1.630057	24.288677	13.723158	80.417675
15.667909	40.094439	12.979045	10.880880	99.892452	21.825691
51.293239	83.911223	61.263983	29.603162	63.755227	52.428719

The outputs:

rank 2 in node051: 4 – 4

rank 0 in node051: 1 – 2

rank 1 in node051: 3 – 3

Round 1 starts:

rank 2 finished averaging:

84.018772	39.438293	78.309922	79.844003	91.164736	19.755137
33.522276	76.822959	27.777471	55.396996	47.739705	62.887092
36.478447	51.340091	95.222973	91.619507	63.571173	71.729693
14.160256	60.696888	1.630057	24.288677	13.723158	80.417675
15.667909	43.313766	28.467340	41.690834	27.546239	21.825691
51.293239	83.911223	61.263983	29.603162	63.755227	52.428719

rank 0 finished averaging:

84.018772	39.438293	78.309922	79.844003	91.164736	19.755137
33.522276	38.019533	76.438212	61.745172	68.254999	62.887092
36.478447	67.305317	43.091782	59.619954	56.203016	71.729693
14.160256	60.696888	1.630057	24.288677	13.723158	80.417675
15.667909	40.094439	12.979045	10.880880	99.892452	21.825691
51.293239	83.911223	61.263983	29.603162	63.755227	52.428719

rank 0 finished averaging:

84.018772	39.438293	78.309922	79.844003	91.164736	19.755137
33.522276	76.822959	27.777471	55.396996	47.739705	62.887092
36.478447	51.340091	95.222973	91.619507	63.571173	71.729693
14.160256	26.806211	48.296895	29.463400	67.042494	80.417675
15.667909	40.094439	12.979045	10.880880	99.892452	21.825691
51.293239	83.911223	61.263983	29.603162	63.755227	52.428719

```
Thread 2 is waiting for signal
Thread 1 is waiting for signal
Thread 0 is broadcasting signals
Round 2 starts now:
...
Total rounds: 18
Result :
```

84.018772	39.438293	78.309922	79.844003	91.164736	19.755137
33.522276	44.443627	60.947406	68.739927	72.209024	62.887092
36.478447	42.971455	53.945278	60.512557	67.063370	71.729693
14.160256	38.668366	49.007878	54.970752	62.356437	80.417675
15.667909	47.086171	51.117751	45.666212	48.624188	21.825691
51.293239	83.911223	61.263983	29.603162	63.755227	52.428719

Check final result with non-parallel version, they are the same.

It is very difficult to test more threads or larger dimensions. The output are huge. What I did is random the input array once, output to a file. Then run non-parallel version and this program on the same array, compare the result of each round using the same method which is described in 3.2. The testing results show that my output are all correct.

Chapter 4

Scalability Investigation

I use three square arrays with different dimensions to do the scalability investigation:

1. 100 x 100 array with precision 0.01
2. 200 x 200 array with precision 0.1
3. 200 x 200 array with precision 0.1

4.1 Time spent

Figure 4.1 shows time spent with different threads. In general, using distribute memory increase the time needed drastically.

From Table 4.1, we could see that using processors in different nodes does not effect time spent on program very much. This may be because it uses distribute memory, so even processors in same node need to store messages to their own local memory. Therefore, both processors in same node and in different nodes doing the same thing.

	100x100 - 0.01	100x100 - 0.1	200x200 - 0.1
1 node, 6 processors each	30s	18s	3m 11s
2 node, 3 processors each	37s	17s	3m 6s
3 node, 2 processors each	38s	21s	3m 6s

Table 4.1: Time spent for 6 processors




Figure 4.1: Time spent



Figure 4.2: Speedup

4.2 Speedup

Figure 4.2 shows the parallel algorithm in comparison with the corresponding sequential algorithm. The maximum speedup is 0.125 only. The reason is that the problem to solve is too small compared with the overhead of creating a message, sending, waiting, reading the reply. In my program, broadcasting array and flag ('isEnd'), sending and receiving calculating results are huge overhead comparing with the calculation. This kind of simple problem does not worth to use distributed memory.

4.3 Efficiency

Figure 4.3 shows efficiency of this program run in parallel. In general, the efficiency decreases as the number of processors increase. The highest efficiency is only 0.036. This means we are not even using 3.6% of each processor. This is a huge waste.

4.4 Compare with shared memory

From Figure 4.4 we can see that even through the overall speedup and efficiency are not good, program using shared memory are a lot better than the one using distributed memory.




Figure 4.3: Efficiency




Figure 4.4: Time spent on 200x200 array with 0.1 precision